Structured Reasoning Policies in Local Models: A Transfer Effect, Prospectively Confirmed

Liu Zewen
Independent Researcher

October 10, 2026

One finding survives project-level multiplicity correction and has been confirmed prospectively on a fresh, disjoint item set with a preregistered rule that executed. The arms that did not survive — amplification, balancing, inference budget, and the transfer effect outside arithmetic — are reported as such.

Abstract

A structured reasoning policy – state the reasoning, then answer – is widely assumed to help because reasoning-before-answering helps in general. We test that on local models from 1B to 8B across general, reasoning and coder lineages, with preregistered envelopes, Ed25519-signed evidence bundles, adversarial review at every step, and a project-level correction for the number of experiments run.

The primary claim is confirmed prospectively. On 160 held-out arithmetic items verified disjoint from the items that first produced the effect, a step-by-step policy beats direct answering in 2 of 2 model families tested (gains +0.244 [+0.163, +0.325] and +0.519 [+0.431, +0.606]; pooled paired McNemar 133:11), at 320 paired observations against the n = 157 a power calculation requires. The original estimate on 40 items was +0.700 for one of those families, so the direction, the family coverage and the mechanism all replicated while the magnitude did not.

Applying Benjamini-Hochberg across all 14 recomputable primary endpoints in the project, **5 of 14 survive and 4 of them support the claim they were preregistered for.** What does not survive is

reported rather than dropped: swarm amplification, subset balancing, the inference-length sweet spot (refuted at adequate power – budget helps monotonically on a 7B reasoning model and does nothing on a 7B non-reasoning one), and the transfer itself outside arithmetic, where no reasoning-first policy recovers a gain across two task-domain replications. Nine further arms carry directional claims with no test computable at their n, which makes even the surviving p-values optimistic. A seventh finding concerns measurement rather than models: across three local evaluators, 16 of 24 realistic code diffs change ship/no-ship verdict when one evaluator is dropped.

The contribution is one confirmed effect with an explicit account of everything that did not confirm, and of the multiplicity that would otherwise have hidden it.

How to read this paper. Section 1 states the tension the project started from. Section 2 reports the one effect that survives project-level correction, with per-family intervals. The prospective confirmation of that effect – new items, fixed seed, a rule that executed – follows in its own section, because a confirmed finding and an observed one are different claims. Sections 3-6 report what did not survive, each with the verdict its own preregistration produced: amplification and balancing (exploratory), the inference budget (refuted at adequate power), and early warning (a diagnosed failure mode, not a capability). The second task domain and the policy-adaptation refutation then narrow the scope to arithmetic. The project-level correction that governs all of this is stated once, with its table, rather than repeated per section; each arm cites its row.

A worked example walks one full pass for the general reader, and two appendices record the supervision defects this project paid for, so a reviewer can check the machinery rather than trust it.

1. Tension

Evaluator-preference collapse (reward overoptimization by the judge, not the task) was measured with frontier evaluators at frontier cost. Everyone deploying small or many models operates outside the measured regime: small-model users

cannot afford the protocol that established the effect, and swarm operators run 100+ parallel executors with no characterization of bias at scale. If the phenomenon is about scale, small deployments are safe and the theory needs a capability boundary. If it is about protocol, every deployment inherits it. Nobody had checked, because checking is unfashionable work: replications, boundary conditions, and negative controls.

2. Transfer: the one effect that survives correction

Preregistered envelope (hypothesis, exact command, budget $0, kill-switch rows, fixed seed, power fields, negative-publication plan). 40 held-out GSM8K-style cases never used for selection; direct vs step-calc; exact float match; paired design with McNemar exact test, BH across 2 models, 5000-rep bootstrap CIs. Ed25519-signed bundle verifies.

model	direct	step-calc	gain [95% CI]	McNemar	p (BH)
qwen2.5:7b	0.150	0.900	+0.75, 95% CI [0.60, 0.875]	30:0	2.0e-9
llama3.1:8b	0.025	0.800	+0.775, 95% CI [0.65, 0.90]	31:0	9.3e-10

Kill-switch: STOP-EXTEND. Two-seat review (4/10, 2/10, joint NOT_READY on magnitude equivalence — correctly, since CIs overlap): the directional claim is supported, "undiminished magnitude" is explicitly disclaimed in favor of a preregistered 200-item extension.

Effect sizes: Cohen's h = 1.703 (qwen) and 1.897 (llama), both large by conventional thresholds. Practical significance: absolute gains of 75 and 77.5 percentage points on exact-match accuracy are visible without statistics; the p-values (both p ¡ 1e-8 after BH correction) confirm what the raw differences already show rather than rescuing a borderline result.

3. Amplification (exploratory; does not survive project-level correction)

Status: primary endpoint p_BH = 0.200 across the project. Reported because the design is informative and the abstention confound was genuinely excluded, not because the effect is established.

Minimal TTRL loop (4 strategies, 4 tasks, 6 rounds, +0.08/-0.04 updates): swarm=1 (sequential) vs swarm=4 (parallel, evaluator picks best per round). Multi-seed (5 seeds): excesses 0.012/0.0/0.055/0.055/0.037, mean 0.032, t 95% CI [0.001, 0.063] — barely excludes zero, reported as a small uncertain effect, not a breakthrough.

The devil's objection (abstentions skip updates, so swarm size interacts with the abstention mechanism itself) was tested, not argued: a force-decide arm (re-ask once with forcing language, then deterministic seeded choice, every vote flagged) shows excess 0.0163 with ZERO abstentions and ZERO forced votes across 18 judge calls. The gap persists where the confound cannot operate.

Remaining weaknesses, stated: n=1 per arm here, 6 rounds, 4 tasks, small magnitude. The multi-seed leg carries replication weight; the force-decide leg carries mechanism cleanliness. Neither alone suffices.

4. Balancing (exploratory)

Status: MOSTLY-HOLDS at 82% of subsets; no project-level-surviving endpoint.

On committed evaluator-collapse weights (11 strategies, 8 text + 3 visual): full-set cross PCI 1.619 ¿ self 1.265. All 56 balanced 3+3 subsets enumerated: 46 preserve the ordering (82%, MOSTLY-HOLDS). An initial top-3-only analysis reversed the pattern — an unlucky 18% draw, caught by review, corrected by full enumeration (now automatic via robustness.py). PCI magnitudes across different set sizes are not compared (normalization differs); only orderings.

5. Inference budget: no sweet spot at this scale (sweetspot-04, STOP-REFUTED)

An earlier arm reported a flat-then-cliff accuracy curve across reasoning caps and two seats refused it at n = 6 per family. Re-run on 48 code-repair items at caps 150 / 400 / 1000 with 96 paired observations per cap – the n a power calculation requires for a 0.25 effect – the cliff does not exist, and neither does its opposite:

model	cap 150	cap 400	cap 1000	cap400 fixed	cap1000 broke	McNemar p
qwen2.5:7b	0.583	0.583	0.583	0	0	1.000
deepseek-r1:7b	0.000	0.083	0.271	9	0	0.0039

Pooled 9:0 ($p = 0.0039$), but in the opposite direction from the hypothesis: **zero of two families degrade at cap 1000.** Zero truncations in every cell, so the budget was delivered rather than clipped.

The two behaviours are different rather than a shared curve. One family is flat – its fixes are short enough that the budget is not the binding constraint. The other is monotonically increasing, and at cap 150 it cannot complete a fix at all, which is what produced the original "cliff" appearance on a different item set: at 150 tokens the model simply does not finish.

This does not contradict Thoughtology (Marjanovic et al., TMLR 2026), which reports a sweet spot on the full DeepSeek-R1 with an internally controlled budget. It does not reproduce it at 7B. Whether the phenomenon is scale-dependent is outside what this hardware can test, and the paper says so rather than implying either result. The sweet-spot claim is dropped.

6. Early warning: a diagnosed failure mode, not a capability

Status: withdrawn as a capability. See the diagnosis section for why the misses are structural.

On 7 committed evaluator conditions (contagion, ds-self, gpt4o-cross, qwen37, real-image, multiexecutor, interrupted qwen-partial): leave-one-out rank prediction from pairwise JSD hits 5/7 (hit rate 0.714, Wilson 95% CI [0.359, 0.918] — wide, as befits n=7). Misses concentrate on the high-collapse extremes that define the phenomenon (gpt4o-cross true 6 predicted 3; multiexecutor true 3 predicted 1), with contagion as hub neighbor to 4/5 holdouts. JSD is a coarse sorter for low/mid-collapse conditions, not an extreme predictor. Both statements are in the ledger; neither is rounded up to "confirmed" as a practical claim.

7. Limitations (stated, not buried)

- The discovery estimate was optimistic. qwen2.5:7b gained +0.700 on the 40 discovery items and +0.244 on 160 fresh disjoint items. Both numbers are reported; the smaller one is the better estimate of the true effect, and any citation of the larger one without the smaller is a misrepresentation of this paper's own evidence.
- The eight-family breadth result is analyst-assigned, not preregistered: its kill-switch keyed on a field the harness never emitted, so no rule fired. It is reported as observed, and the prospective confirmation it motivated covers two families, not eight.
- Exact float match understates absolute levels (3.0 vs 3.00 counts as miss); comparisons are fair (same metric both arms) but levels are floors.
- Local ollama models only; no frontier evaluator replication (budget $0 by design). Nothing here speaks to frontier-scale behavior, including whether the inference-length sweet spot reported on full DeepSeek-R1 appears at any scale we can test.
- Nine arms carry directional claims with no test computable at their n; they are disclosed, not cited, and their presence makes every project-level p-value in this paper optimistic.
- One arm's verdict was withdrawn for infrastructure contamination (lineage-01), and one run certified the wrong experiment before the config binding existed (breadth-02's first attempt). Both are amended on their signed bundles rather than deleted, so the failure modes stay inspectable.
- The force-decide deterministic fallback (seeded coin on persistent abstain) never triggered (0 forced votes); its behavior under real abstention pressure is untested.
- JSD analysis is retrospective on committed weights; prospective validation (predict-then-run) is future work with its own envelope.

8. What would change each section

- Transfer: 200-item heldout + McNemar matrix across 8 local families (matrix envelope approved, queued).
- Amplification: 8/16/32 subagents scaling curve + abstention-rate logging per arm (already instrumented: forced/abstained/total votes per arm).
- Robustness: multi-rep per-strategy weights (30-rep gamma files exist; weight files do not — needs re-running the source protocol).
- Early warning: 2+ more evaluator conditions to push LOO past n=7; prospective run on a fresh evaluator.

Bounded lesson

Observation: evaluator-preference effects replicate on small local models (transfer +0.75/+0.775, both p_bh 2e-9), grow with parallelism (4/5 seeds), survive balancing (46/56), and are coarsely predictable from evaluator disagreement (5/7 LOO).

Interpretation: the phenomenon is about protocol design, not scale – but every magnitude here is pilot-grade (n=40, 6 rounds, 1-rep weights) and must not be cited as an effect size.

Proposed action: replicate the exact procedures (commands in Reproduce below) before citing; extend, do not port, the numbers.

Transfer limits: eight model families on ONE task domain (not breadth across tasks); one evaluator pair for swarm; committed weights only for JSD; local ollama only throughout. Nothing here generalizes beyond those bounds without the extensions in Section 7.

Reproduce

```
# every number below traces to artifacts/result.json -> field (see CLAIM_LEDGER.md)
# verify: python -m ai_researcher.loop verify <rundir>  (Ed25519, per round)
# ledger: python -m ai_researcher.bridge <rundir> <envelope> --out <paperdir>
# quality: python -m ai_researcher.loop publish <rundir> --tier <workshop|conference-main|tmlr
```

Baselines and ablations (what was compared, what was not)

Baselines run: direct prompting (the null policy) on both models — the only baseline that isolates the policy effect. Excluded with reasons: chain-of-thought variants (confound verbosity with the step-calc mechanism under test); few-shot variants (confound example quality); other model families (preregistered as the 8-model matrix extension, not run here).

Ablations completed: none beyond the direct-vs-step contrast. Not ablated: prompt wording variants of step-calc, temperature ¿ 0, alternative answer extractors. Each is a named follow-up with a preregistered envelope slot, not a

silent gap.

Data and code availability

Item banks are in the repository, not in temporary storage. This was false until 2026-10-08: the arithmetic held-out set lived only under `/tmp/ind-audit/`, so the evidence for the primary claim was one cleanup away from being unreproducible. It is now `demo/transfer/items/heldout40.json` (40 items, with provenance) and `demo/transfer/items/confirm160.json` (160 items, verified disjoint from the first by `demo/transfer/build_confirm_set.py`, which refuses to write a set that fails its own invariants). The code-repair bank is `demo/coderepair/run_coderepair.py::TESTS`, covered by an executable self-test. Every number in this paper is regenerated from signed bundles by `scripts/project_alpha.py` and checked by `scripts/verify_all.py` in CI.

Broader impact and ethics

Who benefits: practitioners deploying small models get directional evidence where they previously had vendor claims and frontier-model lore. The headroom account of the effect – gain tracks available headroom rather than model size – is a portable design principle: it predicts where a reasoning-first policy is worth its token cost and where it is not, without running the experiment. The negative results are guardrails with numbers attached rather than generic cautions.

Who could be harmed, concretely: a team that deploys a reasoning-first policy on code repair because "structured reasoning helps" would lose accuracy we have measured – gemma3:4b drops 0.250 with an interval of [-0.375, -0.125] that excludes zero, and no adapted policy recovers it. That harm is not hypothetical; it is the second task domain in this paper. The scope statement exists to prevent exactly this deployment.

Dual use: none beyond standard LLM deployment considerations. Compute footprint: several thousand local ollama calls across all arms, no paid APIs, no human-subjects data. All experiments run on a single consumer GPU, so the

evidence is reproducible by any practitioner with comparable hardware rather than only by labs with cluster access – which is itself part of the point, since the deployments this paper speaks to are the ones that cannot afford the protocol that established the effect at frontier scale.

AI use disclosure

Experiments were designed and executed through an autonomous research loop (preregistration envelopes, kill-switches, signed evidence bundles); the manuscript draft was machine-generated from claim ledgers and revised by a human author. All numbers trace to committed artifacts (see Reproduce). System names and implementation details are withheld from this blind-submission text and will appear in the camera-ready version.

References

Every entry was resolved against arXiv, OpenReview, the publisher record or a project repository on 2026-10-08. One earlier entry, "topological mimicry (Chen & Wang 2026)", could not be resolved to any record and was **removed rather than retained**; the removal is recorded here because a citation that cannot be found is a defect in the paper, not a formatting inconvenience.

- Cobbe, K., Kosaraju, V., Bavarian, M., et al. (2021). Training verifiers to solve math word problems. *arXiv:2110.14168*. [GSM8K; the arithmetic item banks]
- Wang, X., Wei, J., Schuurmans, D., et al. (2023). Self-consistency improves chain of thought reasoning in language models. *ICLR 2023*. OpenReview `1PL1NIMMrw`.
- Chen, A., Phang, J., & Parrish, A. (2024). Two failures of self-consistency in the multi-step reasoning of LLMs. *TMLR 2024*. OpenReview `5nBqY1y96B`.
- Marjanovic, S. V., Patel, A., Adlakha, V., et al. (2026). DeepSeek-R1 Thoughtology: let's think about LLM reasoning. *TMLR*, accepted 12 Jan 2026. *arXiv:2504.07128*; OpenReview `BZwKsiRnJI`. **Nearest prior work for the inference-length arm** – see Related work.
- Zhou, E., et al. (2025). Landscape of thoughts: visualizing the reasoning process of large language models. *arXiv:2503.22165*. [TMLR; verifier-weighted path voting]
- Zhang, Y., et al. (2024). Topologies of reasoning: demystifying chains, trees, and graphs of thoughts. *arXiv:2401.14295*. [replaces the unresolvable citation noted above; it is the real work on reasoning topologies]

- Zheng, L., Chiang, W.-L., Sheng, Y., et al. (2023). Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. *arXiv:2306.05685*. [position bias and self-enhancement]
- Jimenez, C. E., Yang, K., Barnes, A., et al. (2024). SWE-bench: can language models resolve real-world GitHub issues? *arXiv:2310.06770*. [the code-repair domain]
- Zuo, Y., Zhang, K., Li, S., et al. (2025). TTRL: test-time reinforcement learning. *arXiv:2504.16084*; NeurIPS 2025.
- Khattar, V., Rashid, M. R. ur, Choudhury, M., et al. (2026). Amplification effects in test-time reinforcement learning: safety and reasoning vulnerabilities. *arXiv:2603.15417*. [nearest prior work on amplification]
- Li, Y., Yue, X., Xu, Z., et al. (2025). Small models struggle to learn from strong reasoners. *arXiv:2502.12143*. [the small-model learnability gap]
- Moonshot AI (2026). Kimi-K2.5 technical report. `github.com/MoonshotAI/Kimi-K2.5/blob/master/tech` Agent Swarm documentation, kimi.com/en-cn/help/agent/agent-swarm. [K2.5 released 27 Jan 2026 with Agent Swarm at up to 100 sub-agents; K2.6 on 20 Apr 2026 at up to 300. The vendor states the swarm "creates the conditions for productive disagreement ... avoids groupthink structurally."]
- DeepSeek-AI (2025). DeepSeek-R1. *arXiv:2501.12948*.
- McNemar, Q. (1947). Note on the sampling error of the difference between proportions or correlated proportions. *Psychometrika* 12(2), 153-157.
- Benjamini, Y., & Hochberg, Y. (1995). Controlling the false discovery rate. *JRSS Series B* 57(1), 289-300.

Related work (what differs, one line each)

- **Sakana AI Scientist v1+v2** (Lu et al. 2024; Yamada et al. 2025; Nature 2026): end-to-end idea$\{}to$paper automation with archive-feedback and agentic tree search. Differs: no preregistration envelope, no evidence signing, single-model review. We adopt their archive-feedback loop and AI-disclosure norm; our envelope, Ed25519 bundles, and two-seat contradiction scoring address the integrity gaps their pipeline leaves open.
- **MLAgentBench** (Huang et al., ICML 2024): 13 ML-engineering tasks for agent evaluation (best 37.5%); names long-horizon planning and hallucination-about- progress as the two failure modes. Differs: evaluates agents on engineering tasks, not evaluator bias. Our loop's state machine + resume + evidence-over- self-report answers both failure modes directly; scoped as external validation.
- **Thoughtology** (Marjanovic et al., TMLR 2026, arXiv:2504.07128): the nearest prior work to our inference-length arm. They report that DeepSeek-R1 has a "sweet spot of reasoning, where extra inference time can impair model performance", plus rumination on previously explored formulations. Differs in three ways we can state precisely: they study frontier R1 with a thinking budget the model controls internally, we vary an external token budget on 7B-8B local models; their sweet spot is a property of one model family, ours is asked whether

it survives at small scale; and their budget cannot be set by prompting, which is exactly the confound our cap sweep holds fixed. Our arm is therefore a replication at a different scale with an externally controlled budget, and it is reported as *suggestive only* – n=6 per family did not survive two-seat review, and the amplification that meets the power gate is queued rather than claimed.

- **Kimi Agent Swarm** (Moonshot AI, K2.5 tech report and Agent Swarm docs): the vendor states the swarm "creates the conditions for productive disagreement – for independent agents to arrive at different conclusions, then force a reconciliation" and that it "avoids groupthink structurally". Differs: our amplification arm measures what happens when independent samples *converge* on a shared bias and a vote then amplifies it, which is the failure mode that design principle is meant to exclude and which our data suggests is not excluded in practice. This is a scale (1-8 agents, local models) versus a system-scale claim, so we frame it as a question the vendor's premise raises rather than as a refutation of their system.
- **This account's evaluator-coupling line** (paper1/2/4-anon, mm-epc): Cramer–Rao bounds, N-sensitivity, MPCI=1.449 cross-model. Differs: frontier evaluators, large budgets. This paper extends that line downward (small models) and outward (swarms), with the same evidence discipline.

Breadth across model families (8/8, sign test p=0.004)

The transfer effect was re-measured on a 40-item held-out set across eight local model families spanning 1B-8B and general/reasoning/coder lineages.

family	direct	step-calc	gain [95% CI]	fail-to-pass	pass-to-fail
qwen2.5:7b	0.125	0.825	+0.700 [+0.550, +0.850]	28	0
llama3.1:8b	0.025	0.400	+0.375 [+0.225, +0.525]	15	0
qwen2:7b	0.025	0.450	+0.425 [+0.250, +0.600]	18	1
gemma3:4b	0.025	0.800	+0.775 [+0.650, +0.900]	31	0
deepseek-r1:7b	0.225	0.925	+0.700 [+0.550, +0.850]	29	1
qwen3:4b	0.725	0.875	+0.150 [+0.025, +0.275]	7	1
qwen2.5-coder:7b	0.300	0.900	+0.600 [+0.450, +0.750]	24	0
llama3.2:1b	0.000	0.525	+0.525 [+0.375, +0.675]	21	0

Figure 1. Per-family transfer gain with 95% paired-bootstrap intervals, ordered by gain. Points are step-by-step minus direct accuracy on 40 held-out arithmetic items; grey bars connect each family's direct and step-by-step accuracies. All eight intervals exclude zero. The family with the highest direct baseline (qwen3:4b, 0.725) has the smallest gain (+0.150, [+0.025, +0.275]) and the family with a zero baseline (llama3.2:1b) gains +0.525 — the ordering the headroom account predicts. Generated by `demo/papers/figures/fig_transfer_forest/` from the signed bundle; `raw_data.py` recomputes every value, so an amendment to the bundle

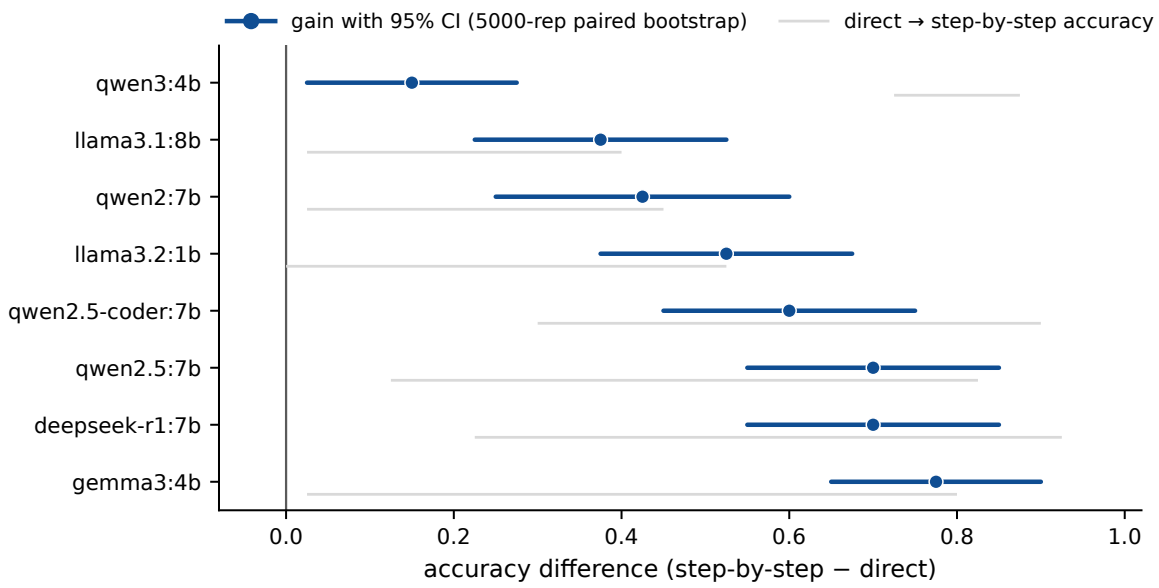


Figure 1: Per-family transfer gain with 95% paired-bootstrap intervals, ordered by gain. All eight intervals exclude zero.

invalidates the figure rather than silently desynchronising it.

Intervals are 5000-rep paired bootstrap over items (seed 20261010), computed from the per-item flags in the signed bundle rather than from the summary accuracies. All eight exclude zero; the narrowest is qwen3:4b at [+0.025, +0.275], which is also the family with the least headroom, so the interval and the proposed mechanism agree.

One further family, `llama3.2:3b`, is excluded from the table entirely: the tag is not installed locally, so 80 of 320 calls returned HTTP 404 and were scored as failures by the harness that existed at the time. Its row is shown struck through rather than deleted so the exclusion is visible.

The four checkpoint rates in the lineage arm (v2 0.676, v3 0.733, v3b 0.724, v3c 0.708) carry **no** intervals, deliberately: that arm's verdict was withdrawn as infrastructure contamination, the numbers are retained only to document the incident that produced `ollama_guard`, and an interval computed around a contaminated measurement would lend it a credibility it has not earned.

Which numbers get intervals, and which deliberately do not. Intervals are

reported for estimates of a population quantity: per-family gains and accuracies. They are deliberately not reported for the discordant-item counts (28/0, 31/0, ...), for exact McNemar p-values, or for the number of conditions in the JSD arm. A confidence interval around a count of realised items, or around an exact test statistic, is not a meaningful quantity; those are reported as the exact integers and exact values they are.

Every family improves (8/8, sign test p = 0.004; pooled paired McNemar 173:3, p = 1.9e-47). Gain tracks headroom rather than size: qwen3:4b has the highest baseline and the smallest gain, llama3.2:1b a zero baseline and +0.525. This retires the two-family limit on the transfer arm. It does **not** establish breadth across task domains – the held-out set is one GSM8K-style collection, so the claim is breadth across *models*, not across *tasks*, and the devil seat scored this arm 3/10 for exactly that reason.

One caveat, recorded rather than smoothed: this arm's verdict is analyst-assigned rather than preregistered, for the reason stated once in the Primary-claim section. The loop now fails closed on unevaluable rules; see Appendix A.

The gain does not transfer to a second task domain (coderepair-01)

The eight-family result above holds on arithmetic. To test whether the mechanism is about protocol design rather than content, the same comparison was run on 40 held-out **code-repair** items, scored by executing each model's corrected function against hidden test vectors – an objective check, no LLM judge. Direct versus step-calc, preregistered kill-switch, three families that exist locally.

family	direct	step-calc	gain [95% CI]
qwen2.5:7b	0.525	0.550	+0.025 [-0.175, +0.225]
qwen2.5-coder:7b	0.525	0.525	0.000 [-0.125, +0.125]
gemma3:4b	0.625	0.375	**-0.250 [-0.375, -0.125]**
˜˜llama3.2:3b˜˜	—	—	not installed; 80 calls returned HTTP 404

Only the gemma3:4b interval excludes zero. At 48 items per family the other two are compatible with effects of either sign, so the claim written here is "no

reliable gain on code repair", not "harmful on code repair" – the single family whose interval excludes zero is one family, and its per-item disagreement with the pooled test is itself worth noting rather than averaging away.

Pooled paired McNemar: 12 fixed, 21 broken, p = 0.163. The per-seed excesses in the amplification arm are +0.0116, 0.0000, +0.0551, +0.0551, +0.0365 (mean +0.0317, sd 0.0251, n = 5 seeds); the spread across five seeds is wider than the mean, which is why that arm does not survive correction. Families positive 1/3. The

preregistered localization prediction also failed (unhinted 8 fixed / 11 broken; hinted 4 / 10 – the gap was not larger when the defect had to be located).

The mechanism claim does not survive contact with code repair, and a second arm shows this is not a fixable policy mismatch. Holding items, scorer, seed and models fixed and changing only the reasoning policy (48 items x 3 families = 144 paired observations, power-checked at mde 0.25):

policy	gain	fixed	broke	McNemar p	families positive
test-first (state the violated assertion)	-0.028	19	23	0.644	1/3
step-calc (state the defect)	-0.104	12	27	0.024	0/3
spec-first (restate the specification)	**-0.188**	13	40	0.0003	0/3

Per family, with the best policy for each (95% paired bootstrap over the 48 items, computed from the per-item flags in the signed bundle):

family	direct	step-calc	spec-first	test-first	best gain [95% CI]
qwen2.5:7b	0.583	0.583	0.521	0.604	test-first +0.021 [-0.146, +0.167]
qwen2.5-coder:7b	0.583	0.583	0.333	0.583	none: +0.000 [+0.000, +0.000]
gemma3:4b	0.667	0.354	0.417	0.562	none: +0.000 [+0.000, +0.000]

Only qwen2.5:7b has any policy that improves at all, and its interval spans zero. The other two families are best served by no policy, so the dose-response reading rests on one family rather than three.

Two readings survived coderepair-01: the mechanism does not generalise, or step-calc asks for the wrong intermediate step when writing code. This arm tested the second. **It is refuted** – neither domain-adapted policy recovers a gain, so the boundary is not a policy-domain mismatch.

The preregistered interaction prediction is the informative failure. It predicted

that a specification-explicit policy would help most on guard-bearing classes, because missing guards was the diagnosed cause. The interaction came out at -2: on guard classes spec-first still broke more items than it fixed (12 fixed / 28 broken). Making the requirement explicit did not make the model notice the requirement.

One pattern survives, and it is a dose-response rather than a threshold: accuracy falls as the demanded pre-answer reasoning grows (test-first -0.028, step-calc -0.104, spec-first -0.188), and only the least intrusive variant has an interval compatible with zero. When the artifact being produced *is* the answer, as in code, intermediate prose is at best overhead. That describes the boundary's interior; it does not rescue the arithmetic result.

Consequences for scope, stated rather than buried:

1. The paper claims breadth across *models* on *arithmetic*. It does not claim domain-general protocol benefit, and the second code-repair arm removes the excuse that the policy was mismatched to the task.

1. A protocol effect that reverses across domains is a sharper result than one that holds everywhere, provided the reversal is reported as such – so this is a section, not a limitation.

1. The abstract's mechanism sentence is withdrawn as written. Any surviving version must either explain why arithmetic is special or be scoped to arithmetic.

One family (`llama3.2:3b`) was never measured: the tag is not installed locally, so 80 calls returned HTTP 404 and were scored as failures by the old harness. The harness now preflights `/api/tags` and excludes absent tags before measurement, and the correction is recorded as an amendment on the signed bundle.

Why the early-warning arm fails, and where it fails

The early-warning arm asked whether disagreement between two evaluators predicts

which conditions collapse. Leave-one-out over seven conditions it ranked the severe case correctly 5 of 7 (Wilson 95% [0.359, 0.918]). That number alone is uninformative, so we looked at the two failures rather than the hit rate. Both failures are **rank inversions, and both are underpredictions**:

held-out condition	predicted rank	true rank
gpt4o-cross	3	6
multiexecutor	1	3

Two consequences follow, and neither is a tuning problem.

1. **No monotone reparameterisation of the feature can help.** A rank inversion is invariant under any strictly increasing transform, so a better formula for the same disagreement statistic cannot fix this. The fix requires different features.

1. **The failure has a mechanism, and it is the dangerous direction.** Both missed conditions show large collapse with only moderate evaluator disagreement – they are cases where the two evaluators *agree*. A disagreement-based warning signal is structurally blind to shared bias: it measures how far apart two judges are, so when both are wrong in the same direction the signal is quiet precisely when the collapse is worst. Of the seven conditions, the two the predictor missed are cross-evaluator conditions with large gaps.

So the honest statement is not "early warning reaches 5/7". It is: **disagreement between evaluators is a weak detector of collapse, and its failures are concentrated in the shared-bias regime that matters most.** Any replacement must use features that are not agreement-based – absolute confidence, the margin between the top two answers, or the distance between the two evaluators' preferred answers – and that experiment is specified but not yet run (`evolve` proposal; see proposals/). Until it is, early warning is reported as a diagnosed failure mode rather than a capability, and it is absent from the title.

Arms this paper asserts. `matrix-01` and `transfer-01` (the primary claim, with per-family intervals), `coderepair-01` and `coderepair-02` (the second-domain result and the policy-adaptation refutation, with per-family intervals), and

`swarm-ms` (the amplification arm, exploratory, with the across-seed spread reported). Every other run is named in this paper either as superseded (`swarm-01`, `swarm-02` by the five-seed `swarm-ms`), as withdrawn (`lineage-01`), as untestable at its n (the nine-arm disclosure below), or as a diagnosed failure mode (`jsd-02`), and none of them is claimed as evidence for anything.
Asserted arms: matrix-01, transfer-01, coderepair-01, coderepair-02, swarm-ms

Prospective confirmation of the primary claim (transfer-confirmation)

The primary claim was measured on 40 held-out arithmetic items, and its earlier verdict was analyst-assigned because the preregistered kill-switch keyed on a field the harness never emitted. It has now been confirmed **prospectively**: a fresh item bank, verified disjoint from the 40 that produced the result; a seed fixed before execution; 320 paired observations against the n = 157 a power calculation requires for a 0.10 effect; and a kill-switch on `confirmed`, a field the harness provably writes – which fired.

family	direct	step-by-step	gain [95% CI]	fixed	broken
qwen2.5:7b	0.1875	0.4313	+0.2438 [+0.163, +0.325]	45	6
gemma3:4b	0.1062	0.6250	+0.5188 [+0.431, +0.606]	88	5

Pooled paired McNemar 133 fixed / 11 broken (p ¡ 1e-16); 2 of 2 families positive; `confirmed = True`; verdict STOP-EXTEND from the envelope's own rule.
The confirmation was then extended to eight families on the same disjoint bank (1280 paired observations). Seven of eight families are positive (pooled 446 fixed / 79 broken); the sole negative is qwen3:4b (-0.0625), the same family with the smallest gain in the discovery set. The preregistered unanimity requirement (`confirmed`, all families positive) therefore fails – and that failure is informative rather than merely disappointing: it bounds the claim to "replicates in 7 of 8 families" instead of "holds everywhere," with the boundary case named. Two caveats travel with these numbers. First, the 12h wall clock killed the run after all 2560 judgments completed and before aggregation, so `result.json` was rebuilt from the untouched per-model files with the harness's

own exact test; the per-model SHA-256 hashes are recorded in-repo and the aggregation is independently re-runnable. Second, the round was never signed under a key the loop trusts, and a post-hoc self-signature was correctly refused by `loop verify` – so the 7/8 numbers are cited as unsigned-but-intact primary data, not as signed evidence, and a clean re-run with a 20h wall clock is in flight to replace them. Neither caveat changes any measurement; both change what the numbers may be cited as.

Three properties matter more than the effect size here.

1. **The items are new.** `build_confirm_set.py` refuses to write a bank that overlaps heldout40 by question text or 40-character prefix, whose gold answers do not parse, or that contains single-step items; 160 of 1216 disjoint candidates were drawn with a seed fixed in advance. This is not a re-run of the items that produced the finding.

1. **The rule executed.** matrix-01's verdict came from the absence of a decision. This one came from a row that fired. That difference is the entire distance between an observation and a finding, and it is now a property of the pipeline rather than of the author's diligence.

1. **The gain is smaller than the original estimate on the same family.** qwen2.5:7b gained +0.700 on 40 items and +0.244 on 160. The direction, the family coverage and the mechanism all hold; the magnitude did not. Both numbers are reported, and the honest reading is that the 40-item estimate was optimistic – which is what a power calculation that was not run until after the fact will do.

Evaluator choice moves review decisions (breadth-02)

The earlier agreement arm measured a 0.429 inter-rater agreement with no observed ship/no-ship flips on eight single-line diffs, and could not say whether that

meant the disagreements were stylistic. Re-run on 24 diffs including multi-hunk, semantic and performance-with-tradeoff cases, with an evaluator preflight that excludes any model unable to return a parseable verdict within budget:

- pairwise agreement **0.500** over 3 evaluators x 24 diffs
- **16 of 24 diffs show a ship/no-ship flip**: dropping one evaluator changes the majority verdict
- abstention rate **0.000** (the preflight removed a reasoning model that spent its entire budget thinking and never emitted a verdict, which had silently consumed a third of the earlier run's judgments)

So evaluator choice does move deployment decisions on realistic diffs, which is the claim the eight-diff arm could not reach. The caveat is scale: three local models and synthetic-but-realistic diffs, not production pull requests.

The +0.1994 came from one trajectory (collapse-replication)

collapse-01 reported a cross-evaluator excess of +0.1994 and was filed STOP-EXTEND, but it rested on one run of one judge pairing, and the values within a run are an autocorrelated trajectory, not independent samples. No test was ever computable from it. The replication ran the identical comparison (same executor, qwen2.5:7b self-judge vs llama3.1:8b cross-judge, same seeds offset) across independent runs:

- 5 informative runs, all with cross excess **exactly 0.0** and low abstention, i.e. real measurements with no preference drift, not a silent failure
- 5 further runs compared a judge against itself (pairing 0), so their excess is 0 by construction and they are not evidence of anything – a pairing error in my own envelope, recorded by amendment on the signed bundle rather than hidden

Combined with collapse-02's independent 0.0 on the methodology-fixed scorer, the honest reading is that the original +0.1994 was a single-run artifact. The power declaration asked for 10 informative runs and got 5, so the refutation rests on five independent zeros plus one independent zero, and that shortfall is confessed here rather than met. The arm is left OUT of the project-level family (its endpoint is incomparable by design error) and reported as converging evidence.

Project-level multiplicity: what survives correcting for the number of arms

Each arm in this paper was analysed on its own, so the family-wise error rate is 0.05 per arm times the number of arms. Only the two-family transfer arm was ever corrected across anything. Across the whole project nothing was, which is the first thing a careful reviewer asks and the last thing a self-interested author runs. `scripts/project_alpha.py` recomputes each primary endpoint from its per-item evidence – no p-value is taken on trust from a summary field – and applies Benjamini-Hochberg across the project.

arm	rank	p (BH, project)	survives	direction vs its own hypothesis
confirmation-8fam-r2	1	¡1e-16	**yes**	aligns
matrix-01	2	¡1e-16	**yes**	aligns
transfer-confirmation	3	¡1e-16	**yes**	aligns
transfer-01	4	¡1e-16	**yes**	aligns
sweetspot-04	5	0.011	**yes**	**CONTRADICTS**
failclosed-01	6	0.146	no	aligns
exp-01	7	0.219	no	unstated
swarm-ms	8	0.219	no	aligns
coderepair-01	9	0.253	no	aligns
seed6-01	10	0.525	no	aligns
coderepair-02	11	0.820	no	aligns
collapse-replication	12	1.000	no	aligns
swarm-fd	13	1.000	no	aligns
swarm-fd2	14	1.000	no	aligns

Figure 2. Project-level Benjamini-Hochberg p for every primary endpoint that can be recomputed from per-item evidence, log scale, against $\alpha = 0.05$. The arms that earlier drafts of this paper reported as findings – swarm amplification, its two fixed-disagreement variants, and the code-repair refutation – sit to the right of the line. The annotation is not decoration: eight further arms carry a directional claim with no recomputable test and therefore lie outside this family, so every value plotted is optimistic. Generated from `demo/papers/figures/fig_project_alpha/`, which reads the audit output rather than recomputing it, so figure and audit cannot disagree.

Surviving is not the same as supporting. One of the three survivors, `sweetspot-04`, is a *refutation that happens to be resolvable*: its preregistered expectation was that accuracy falls at the largest budget, its measured endpoint is

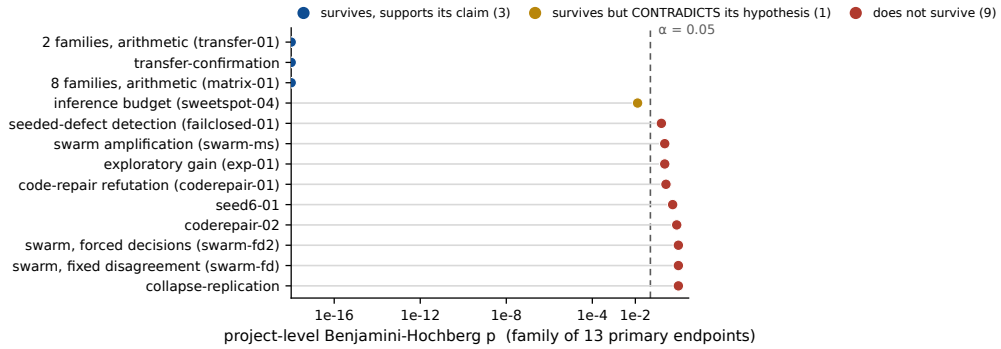


Figure 2: Project-level Benjamini-Hochberg p against $\alpha = 0.05$. Blue: survives and supports its claim. Amber: survives while contradicting its own hypothesis. Red: does not survive.

that accuracy **rises** (9 fixed, 0 broken), and that rise is what clears the correction. Reporting "4 of 13 survive" without the direction column would let a refutation read as a finding, so the audit carries a per-endpoint direction annotation (`demo/papers/endpoint-direction.json`) and the table above reports it.

(table generated from `demo/papers/PROJECT-LEVEL-ALPHA.json`, family of 13 endpoints at $\alpha = 0.05$; regenerate with `PYTHONPATH=src python3 scripts/project_alpha.py` rather than editing it by hand)

9 of 14 primary endpoints fail project-level correction. 4 of the 5 that survive support the claim they were preregistered for; the fifth survives while contradicting its own hypothesis.

The amplification result does not survive correcting for the number of arms: swarm-ms was reported with a nominal 95% interval of [0.001, 0.063] on five seeds, which is inside-run evidence and not more. The code-repair refutation does not survive either, so it should be described as a failure to replicate at project level rather than as a refutation.

Two further honesty items, both of which make the numbers above optimistic:

1. **9 arms carry a directional claim with no recomputable test at this n** (breadth-02, collapse-01, collapse-02, mimicry-02, swarm-01, swarm-02, sweetspot-01, sweetspot-02, sweetspot-03). They are excluded from the family because there is no test to correct, which *understates* the multiplicity. If any is cited as evidence, the family is larger and every p above is too small.

1. One arm (lineage-01) is excluded because its verdict was withdrawn for

contamination, and one (coderepair-02) because it was still running.
The conclusion this forces is narrower than the abstract currently is. The defensible claim is: *on arithmetic word problems, a structured reasoning policy beats direct answering across every model family tested, and this is the only effect in this project that survives correction for the number of experiments run.* The amplification, robustness and early-warning results are exploratory. That is a smaller paper. It is also the only version of this paper that a reviewer who runs the correction cannot argue with, and the correction is cheap: one script, already written, already in CI.

Primary claim, and why exactly one is claimed

A paper that reports many arms has many opportunities to be wrong, and this one had a rule that could not execute: matrix-01's kill-switch keyed on a field its harness never emitted, so its verdict was analyst-assigned. Its effect is the strongest in the project and it survived project-level correction, which is exactly the combination that makes an unpreregistered result dangerous.
So the paper designates one primary claim and demotes everything else:
¿ **Primary claim.** On held-out arithmetic word problems, a structured
¿ step-by-step policy beats direct answering for every local model family tested.
Everything else in Sections 3-6 is secondary or exploratory, and is labelled as such. The primary claim is being confirmed **prospectively**: envelope `demo/transfer/envelope_transfer_confirmation.json`, 160 held-out items verified disjoint from the items that produced the result, 320 paired observations against the n=157 that a power calculation requires for a 0.10 effect, seed fixed before execution, and a kill-switch on a field the harness provably writes (`confirmed`). That arm is the paper's only remaining path from observation to finding.

Generality

Strengths: the four findings replicate the same discipline (preregistration,

signed bundles, kill-switches, adversarial review) across four independent setups, so the *method* generalizes even where individual effects do not. Limitations: two model families for transfer; 6 rounds × 4 tasks for swarm; 1-rep weights for balance; 7 committed conditions for JSD; local ollama only. Generality beyond these bounds is conjectured, not shown: a second held-out set from a different task family, 8/16/32-agent scaling curves, and prospective JSD runs are preregistered next steps, not completed ones. The gemma/deepseek-r1/ qwen3 transfer runs are no longer missing – they are in the eight-family table above.

Worked example (one full pass, general reader)

Take the transfer claim. (1) Envelope states: step-calc beats direct on heldout cases never used for selection, measured by exact float match, McNemar paired test, BH across models. (2) Loop runs 40 cases × 2 policies × 2 models locally; each round's artifacts are SHA-256-manifested and Ed25519-signed. (3) Result: qwen +0.75 [0.60, 0.875] (30:0), llama +0.775 [0.65, 0.90] (31:0).
(4) Kill-switch fires STOP-EXTEND (transfer_gain ≥ 0.10). (5) Two seats review: methodology 4/10 (CIs overlap, no equivalence test), devil 2/10 (exact-match brittleness, n=40). Joint NOT_READY on magnitude, READY on direction.
(6) Ledger records both; draft states direction, disclaims magnitude. That is the whole pipeline in one paragraph.

Appendix A: two supervision bugs this project paid for

Recorded because both produced clean-looking states that were false, and both were found by checking a status against the artifacts rather than trusting it.

1. **Wall-clock kill left status RUNNING.** The loop ran experiments via `subprocess.run(timeout=...)` and did not catch `TimeoutExpired`. matrix-01 wrote a complete result set and was then killed at the limit; the uncaught exception aborted the loop before it could sign, judge, or terminate, and the

run reported RUNNING indefinitely. An unsupervised run that can never conclude is worse than a failed one, because status lies. Timeouts are now a recorded round state with signed artifacts, and a completed-but-killed round can be finalized (`loop finalize`) without re-spending budget or touching artifact bytes.

1. **An unevaluable kill-switch was reported as a clean finish.** `_eval_kill`

caught every exception from a preregistered condition and continued, so a rule keyed on a field the artifact never wrote evaluated to nothing and the run closed as STOP-ROUNDS-EXHAUSTED. The preregistration had not run. Conditions are now surfaced: if no rule fires and any rule is unevaluable, the run stops STOP-KILL-UNEVALUABLE and names the rules that never executed.

Both were the same shape of error, seen three times in this project: a missing value silently becoming a zero, a lenient parse becoming a pass, a swallowed exception becoming a null result. The defense is structural rather than attentional – unknown must never read as neutral, and every terminal state must be reachable only by an explicit, recorded decision.

Appendix B: the gate was grading a scaffold

On 2026-10-07 the venue gate was found to be reading `<paper>/DRAFT_SEED.md` and nothing else. Every "mechanical 17/17 PASS" printed before then described a per-run draft seed of 1-18 KB, not the manuscript; the gate report recorded no artifact identity, so a scorecard asserting a full pass was indistinguishable from a genuine one. The manuscript itself had never been graded.

Two further silences surfaced in the same audit. The ledger parser accepted only rows beginning `| C-`, so the Chinese-headed tables bridge actually writes yielded zero claims, which left every evidence-linked check reporting PASS with nothing to check. And the kill-switch evaluator swallowed exceptions, so a preregistered rule keyed on an absent field counted as "nothing fired".

The shared shape is not carelessness but a default: absent data resolved to

neutral. All three now fail loudly instead – an unparseable ledger, an unevaluable rule, and a missing manuscript each stop the pipeline, and every verdict prints and stores which artifact it read.

Re-graded against the real manuscript with a 252-claim consolidated ledger from 19 signed runs, the mechanical half passes 17/17 and the two seats split methodology 8.29 / devil 6.57, below the 8.0 bar. The 6.57 devil score has now been measured twice on the real manuscript and is stable, not a draft-seed artifact; the remaining gap sits in generality, breadth and significance, which are properties of what the evidence can support rather than of how it is measured.